

naharia-qst843-assignment3

April 26, 2023

1 Title: Enhancing Logistic Regression Model for Telco Customer Churn Prediction || Assignment 3

author: @nahariap@bu.edu | Prateek Naharia

Objective: In this assignment, you will revisit the steps from the notebook 11-Classification-Regression/02-Logistic-Regression-Example.ipynb, applying them to the Telco-Customer-Churn dataset. Your goal is to enhance the logistic regression model using knowledge and examples from the previous three lectures. The data is located in data/Telco-Customer-Churn.csv.

1.1 Customer Churn

Also known as customer attrition, or customer turnover is the loss of clients or customers. Customer churn is a critical metric because it is much less expensive to retain existing customers than it is to acquire new ones.

Companies usually make a distinction between voluntary churn and involuntary churn. In most analyses involuntary churn is excluded.

Predictive analytics uses machine learning to predict the likelihood of a customer churning. These models will identify a small subgroup of potential customers that are at a higher risk of abandoning the company. This empowers the company to focus on the subset of the customers who are most likely to churn and through customer retention marketing programs stop some of that to happen.

1.2 Data

Telco Customer Churn

The data was downloaded from IBM Sample Data Sets:
<https://www.ibm.com/communities/analytics/watson-analytics-blog/guide-to-sample-datasets/>

Each row represents a customer, each column contains customer's attributes described as below:

- **customerID:** Customer ID
- **gender:** Customer gender (female, male)
- **SeniorCitizen:** Whether the customer is a senior citizen or not (1, 0)
- **Partner:** Whether the customer has a partner or not (Yes, No)
- **Dependents:** Whether the customer has dependents or not (Yes, No)
- **tenure:** Number of months the customer has stayed with the company
- **PhoneService:** Whether the customer has a phone service or not (Yes, No)
- **MultipleLines:** Whether the customer has multiple lines or not (Yes, No, No phone service)

- **InternetService**: Customer's internet service provider (DSL, Fiber optic, No)
- **OnlineSecurity**: Whether the customer has online security or not (Yes, No, No internet service)
- **OnlineBackup**: Whether the customer has online backup or not (Yes, No, No internet service)
- **DeviceProtection**: Whether the customer has device protection or not (Yes, No, No internet service)
- **TechSupport**: Whether the customer has tech support or not (Yes, No, No internet service)
- **StreamingTV**: Whether the customer has streaming TV or not (Yes, No, No internet service)
- **StreamingMovies**: Whether the customer has streaming movies or not (Yes, No, No internet service)
- **Contract**: The contract term of the customer (Month-to-month, One year, Two year)
- **PaperlessBilling**: Whether the customer has paperless billing or not (Yes, No)
- **PaymentMethod**: The customer's payment method (Electronic check, Mailed check, Bank transfer (automatic), Credit card (automatic))
- **MonthlyCharges**: The amount charged to the customer monthly
- **TotalCharges**: The total amount charged to the customer
- **Churn**: Whether the customer churned or not (Yes or No)

The data set includes information about:

- Customers who left - the column is called **Churn**
- Services that each customer has signed up for - phone, multiple lines, internet, online security, online backup, device protection, tech support, and streaming TV and movies
- Customer account information - how long they've been a customer, contract, payment method, paperless billing, monthly charges, and total charges
- Demographic info about customers - gender, age range, and if they have partners and dependents

1.3 1. Data Preparation and Logistic Regression Model

Importing Libraries

```
[150]: from pyspark.ml import Pipeline
from pyspark.ml.feature import VectorAssembler, StringIndexer
from pyspark.ml.classification import LogisticRegression
from pyspark.ml.tuning import CrossValidator, ParamGridBuilder
from pyspark.ml.evaluation import BinaryClassificationEvaluator
from pyspark.mllib.evaluation import MulticlassMetrics
```

```
[151]: from pyspark.sql import SparkSession

spark = SparkSession.builder \
    .appName("YourAppName") \
    .config("spark.jars.packages", "com.google.cloud.spark:
↳spark-bigquery-with-dependencies_2.12:0.21.1") \
    .getOrCreate()
```

```
[152]: %matplotlib inline
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
from pyspark.sql.functions import *
from pyspark.sql import SparkSession

# the following line gets the bucket name attached to our cluster
bucket = spark._jsc.hadoopConfiguration().get("fs.gs.system.bucket")

# specifying the path to our bucket where the data is located (no need to edit,
↳ this path anymore)
data = "gs://" + bucket + "/data/"
print(data)
```

gs://pnqst843/data/

Importing data:

```
[153]: df = spark.read.format("csv")\
.option("header", "true")\
.option("inferSchema", True)\
.load(data + "Telco-Customer-Churn.csv")\
.coalesce(5)

df = df.drop('customerID') # Dropping customerID
df.cache()
df.show(5)
df.printSchema()
print("This datasets consists of {} rows.".format(df.count()))
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
----+-----+-----+
|gender|SeniorCitizen|Partner|Dependents|tenure|PhoneService| MultipleLines|InternetService|OnlineSecurity|OnlineBackup|DeviceProtection|TechSupport|StreamingTV|StreamingMovies|Contract|PaperlessBilling|PaymentMethod|MonthlyCharges|TotalCharges|Churn|
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
----+-----+-----+
|Female|          0|    Yes|          No|          1|          No|NoPhoneService|
DSL|          No|          Yes|          No|          No|          No|          No|
No|Month_to_month|          Yes|    ElectronicCheck|          29.85|
29.85|    No|
| Male|          0|    No|          No|          34|          Yes|          No|
```



```
[154]: [(c, df.where(col(c).isNull()).count()) for c in df.columns]
```

```
[154]: [('gender', 0),
        ('SeniorCitizen', 0),
        ('Partner', 0),
        ('Dependents', 0),
        ('tenure', 0),
        ('PhoneService', 0),
        ('MultipleLines', 0),
        ('InternetService', 0),
        ('OnlineSecurity', 0),
        ('OnlineBackup', 0),
        ('DeviceProtection', 0),
        ('TechSupport', 0),
        ('StreamingTV', 0),
        ('StreamingMovies', 0),
        ('Contract', 0),
        ('PaperlessBilling', 0),
        ('PaymentMethod', 0),
        ('MonthlyCharges', 0),
        ('TotalCharges', 11),
        ('Churn', 0)]
```

From the above output, there are 11 null values in total charges

```
[155]: pd.options.display.max_columns = None # do not truncate the middle columns
        df.limit(5).toPandas()
```

```
[155]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	Female	0	Yes	No	1	No	
1	Male	0	No	No	34	Yes	
2	Male	0	No	No	2	Yes	
3	Male	0	No	No	45	No	
4	Female	0	No	No	2	Yes	

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	\
0	NoPhoneService	DSL	No	Yes	
1	No	DSL	Yes	No	
2	No	DSL	Yes	Yes	
3	NoPhoneService	DSL	Yes	No	
4	No	FiberOptic	No	No	

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
0	No	No	No	No	Month_to_month	
1	Yes	No	No	No	OneYear	
2	No	No	No	No	Month_to_month	
3	Yes	Yes	No	No	OneYear	

4	No	No	No	No	Month_to_month
---	----	----	----	----	----------------

	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Yes	ElectronicCheck	29.85	29.85	No
1	No	MailedCheck	56.95	1889.50	No
2	Yes	MailedCheck	53.85	108.15	Yes
3	No	BankTransferAutomatic	42.30	1840.75	No
4	Yes	ElectronicCheck	70.70	151.65	Yes

```
[156]: null_totalcharges = df.filter(df['TotalCharges'].isNull())
null_totalcharges.toPandas()
```

```
[156]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	Female	0	Yes	Yes	0	No	
1	Male	0	No	Yes	0	Yes	
2	Female	0	Yes	Yes	0	Yes	
3	Male	0	Yes	Yes	0	Yes	
4	Female	0	Yes	Yes	0	No	
5	Male	0	Yes	Yes	0	Yes	
6	Male	0	Yes	Yes	0	Yes	
7	Female	0	Yes	Yes	0	Yes	
8	Male	0	Yes	Yes	0	Yes	
9	Female	0	Yes	Yes	0	Yes	
10	Male	0	No	Yes	0	Yes	

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	\
0	NoPhoneService	DSL	Yes	No	
1	No	No	NoInternetService	NoInternetService	
2	No	DSL	Yes	Yes	
3	Yes	No	NoInternetService	NoInternetService	
4	NoPhoneService	DSL	Yes	Yes	
5	No	No	NoInternetService	NoInternetService	
6	Yes	No	NoInternetService	NoInternetService	
7	No	No	NoInternetService	NoInternetService	
8	No	No	NoInternetService	NoInternetService	
9	Yes	DSL	No	Yes	
10	Yes	DSL	Yes	Yes	

	DeviceProtection	TechSupport	StreamingTV	\
0	Yes	Yes	Yes	
1	NoInternetService	NoInternetService	NoInternetService	
2	Yes	No	Yes	
3	NoInternetService	NoInternetService	NoInternetService	
4	Yes	Yes	Yes	
5	NoInternetService	NoInternetService	NoInternetService	
6	NoInternetService	NoInternetService	NoInternetService	
7	NoInternetService	NoInternetService	NoInternetService	

8	NoInternetService	NoInternetService	NoInternetService	
9	Yes	Yes	Yes	
10	No	Yes	No	

	StreamingMovies	Contract	PaperlessBilling	PaymentMethod \
0	No	TwoYear	Yes	BankTransferAutomatic
1	NoInternetService	TwoYear	No	MailedCheck
2	Yes	TwoYear	No	MailedCheck
3	NoInternetService	TwoYear	No	MailedCheck
4	No	TwoYear	No	CreditCardAutomatic
5	NoInternetService	TwoYear	No	MailedCheck
6	NoInternetService	TwoYear	No	MailedCheck
7	NoInternetService	TwoYear	No	MailedCheck
8	NoInternetService	OneYear	Yes	MailedCheck
9	No	TwoYear	No	MailedCheck
10	No	TwoYear	Yes	BankTransferAutomatic

	MonthlyCharges	TotalCharges	Churn
0	52.55	NaN	No
1	20.25	NaN	No
2	80.85	NaN	No
3	25.75	NaN	No
4	56.05	NaN	No
5	19.85	NaN	No
6	25.35	NaN	No
7	20.00	NaN	No
8	19.70	NaN	No
9	73.35	NaN	No
10	61.90	NaN	No

```
[157]: df = df.fillna(0)
```

1.3.1 RFormula

Define an RFormula that uses all of the columns as features and call it **supervised**:

```
[158]: from pyspark.ml.feature import RFormula

supervised = RFormula(formula="Churn ~ .")
#supervised.fit(df).transform(df).show(2, False)
```

Fit the RFormula transformer and call it **fittedRF**:

```
[159]: fittedRF = supervised.fit(df)
```

Using **fittedRF** transform our **df** DataFrame. Call this **preparedDF**:

```
[160]: preparedDF = fittedRF.transform(df)
```

Print the first couple of rows of `preparedDF`, with the `truncate` option off:

```
[161]: preparedDF.show(2, truncate=False)
```

[illegible]

only showing top 2 rows

Below we will retrieve the name of the columns used to make our feature vector and store them in a pandas DataFrame:

```
[162]: featureCols = pd.DataFrame(preparedDF.schema["features"] .
    ↪ metadata["ml_attr"] ["attrs"] ["binary"] +
    preparedDF.schema["features"] . metadata["ml_attr"] ["attrs"] ["numeric"] ) .
    ↪ sort_values("idx")

featureCols = featureCols.set_index('idx')
featureCols.head()
```



```
# binary_features = preparedDF.schema["features"].
↳ metadata["ml_attr"]["attrs"]["binary"]
# numeric_features = preparedDF.schema["features"].
↳ metadata["ml_attr"]["attrs"]["numeric"]
# all_features = binary_features + numeric_features
# sorted_features = sorted(all_features, key=lambda x: x['idx'])
# featureCols = pd.DataFrame(sorted_features).set_index('idx')
# featureCols.head()
```

```
[162]:          name
idx
0      gender_Male
1    SeniorCitizen
2      Partner_No
3    Dependents_No
4          tenure
```

Split the transformed data into train and test. Use a 30% split and a seed.

```
[163]: train, test = preparedDF.randomSplit([0.7, 0.3], seed=42)
```

1.3.2 LogisticRegression

Instantiate an instance of LogisticRegression. Call it lr:

```
[164]: from pyspark.ml.classification import LogisticRegression

lr = LogisticRegression()
```

Print the parameters of lr to check the default values used. You can always come back to the cell above and change the default values:

```
[165]: print(lr.explainParams())
```

```
aggregationDepth: suggested depth for treeAggregate (>= 2). (default: 2)
elasticNetParam: the ElasticNet mixing parameter, in range [0, 1]. For alpha =
0, the penalty is an L2 penalty. For alpha = 1, it is an L1 penalty. (default:
0.0)
family: The name of family which is a description of the label distribution to
be used in the model. Supported options: auto, binomial, multinomial (default:
auto)
featuresCol: features column name. (default: features)
fitIntercept: whether to fit an intercept term. (default: True)
labelCol: label column name. (default: label)
lowerBoundsOnCoefficients: The lower bounds on coefficients if fitting under
bound constrained optimization. The bound matrix must be compatible with the
shape (1, number of features) for binomial regression, or (number of classes,
number of features) for multinomial regression. (undefined)
```

lowerBoundsOnIntercepts: The lower bounds on intercepts if fitting under bound constrained optimization. The bounds vector size must be equal with 1 for binomial regression, or the number of classes for multinomial regression. (undefined)

maxBlockSizeInMB: maximum memory in MB for stacking input data into blocks. Data is stacked within partitions. If more than remaining data size in a partition then it is adjusted to the data size. Default 0.0 represents choosing optimal value, depends on specific algorithm. Must be ≥ 0 . (default: 0.0)

maxIter: max number of iterations (≥ 0). (default: 100)

predictionCol: prediction column name. (default: prediction)

probabilityCol: Column name for predicted class conditional probabilities. Note: Not all models output well-calibrated probability estimates! These probabilities should be treated as confidences, not precise probabilities. (default: probability)

rawPredictionCol: raw prediction (a.k.a. confidence) column name. (default: rawPrediction)

regParam: regularization parameter (≥ 0). (default: 0.0)

standardization: whether to standardize the training features before fitting the model. (default: True)

threshold: Threshold in binary classification prediction, in range $[0, 1]$. If threshold and thresholds are both set, they must match.e.g. if threshold is p , then thresholds must be equal to $[1-p, p]$. (default: 0.5)

thresholds: Thresholds in multi-class classification to adjust the probability of predicting each class. Array must have length equal to the number of classes, with values > 0 , excepting that at most one value may be 0. The class with largest value p/t is predicted, where p is the original probability of that class and t is the class's threshold. (undefined)

tol: the convergence tolerance for iterative algorithms (≥ 0). (default: $1e-06$)

upperBoundsOnCoefficients: The upper bounds on coefficients if fitting under bound constrained optimization. The bound matrix must be compatible with the shape (1, number of features) for binomial regression, or (number of classes, number of features) for multinomial regression. (undefined)

upperBoundsOnIntercepts: The upper bounds on intercepts if fitting under bound constrained optimization. The bound vector size must be equal with 1 for binomial regression, or the number of classes for multinomial regression. (undefined)

weightCol: weight column name. If this is not set or empty, we treat all instance weights as 1.0. (undefined)

Fit the model on train and call it lrModel:

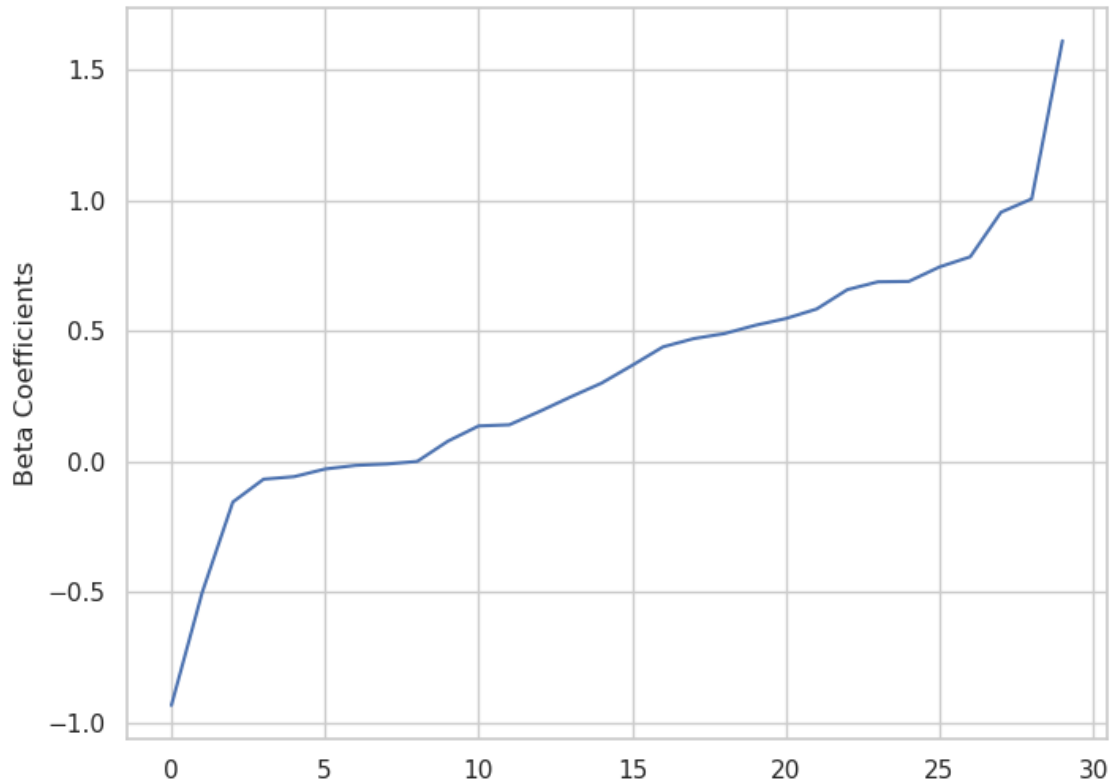
```
[166]: lrModel = lr.fit(train)
```

Below we plot the coefficients of our model in a sorted fashion:

```
[167]: plt.rcParams["figure.figsize"] = (8,6)
beta = np.sort(lrModel.coefficients)
plt.plot(beta)
```

```
plt.ylabel('Beta Coefficients')
```

```
[167]: Text(0, 0.5, 'Beta Coefficients')
```



1.3.3 Feature importance

We already retrieved the name of the features. Let's join it with the coefficients to identify the ones with bigger absolute value:

```
[168]: coefsArray = np.array(lrModel.coef) # convert to np.array
coefsDF = pd.DataFrame(coefsArray, columns=['coefs']) # to pandas

coefsDF = coefsDF.merge(featureCols, left_index=True, right_index=True) # join
↳ it with featureCols we created above
coefsDF.sort_values('coefs', inplace=True) # Sort them
coefsDF.head(10)
```

```
[168]:
```

	coefs	name
23	-0.933980	Contract_TwoYear
9	-0.502347	InternetService_DSL
6	-0.155183	MultipleLines_No
4	-0.067761	tenure

```

28 -0.057943 MonthlyCharges
0 -0.028732 gender_Male
26 -0.014802 PaymentMethod_MailedCheck
2 -0.009636 Partner_No
29 0.000387 TotalCharges
27 0.077996 PaymentMethod_BankTransferAutomatic

```

Plot a Bar Chart

```
[169]: import seaborn as sns
```

```

# plt.rcParams["figure.figsize"] = (20,3)
# plt.xticks(rotation=90)
# plt.bar(coefsDF.name, coefsDF.coefs)
# plt.title('Ranked coefficients from the logistic regression model')
# plt.show()

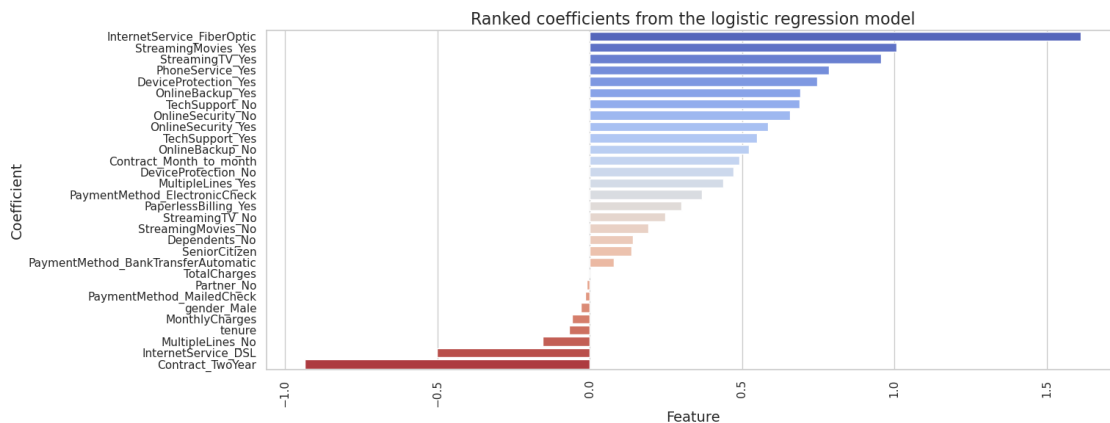
sns.set(style="whitegrid")
plt.rcParams["figure.figsize"] = (15, 6)
plt.rcParams["font.size"] = 12

```

```

[170]: sorted_coefs = coefsDF.sort_values("coefs", ascending=False)
ax = sns.barplot(x="coefs", y="name", data=sorted_coefs, palette="coolwarm")
ax.set_title('Ranked coefficients from the logistic regression model',
             ↪fontsize=16)
ax.set_xlabel("Feature", fontsize=14)
ax.set_ylabel("Coefficient", fontsize=14)
plt.xticks(rotation=90)
plt.show()

```



From our fitted model, `lrModel`, extract the summary and call it `summary`:

```
[171]: summary = lrModel.summary
```

From `summary` extract `areaUnderROC`. Note that this AUC is from the `train` dataset and we should pay more attention to the AUC coming from the `test` set:

```
[172]: print("AUC:", summary.areaUnderROC)
```

AUC: 0.8474526767588296

From `summary` extract `roc` and convert it to a pandas DataFrame. Call it `roc`:

```
[173]: roc = summary.roc.toPandas()
roc
```

```
[173]:
```

	FPR	TPR
0	0.00000	0.000000
1	0.00000	0.003754
2	0.00027	0.006757
3	0.00027	0.010511
4	0.00054	0.013514
...
1000	0.99622	1.000000
1001	0.99757	1.000000
1002	0.99892	1.000000
1003	1.00000	1.000000
1004	1.00000	1.000000

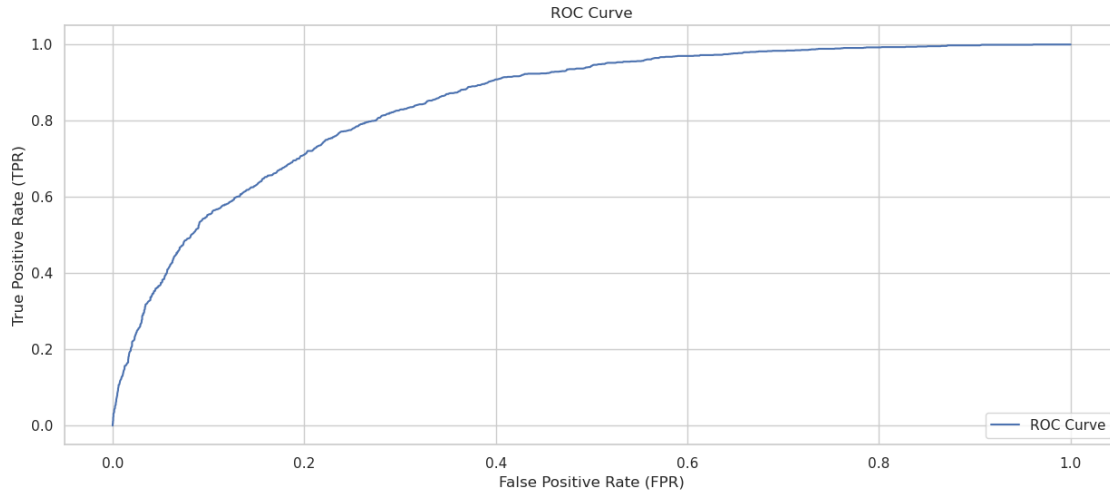
[1005 rows x 2 columns]

Visualize the roc DataFrame:

```
[174]: %matplotlib inline
import matplotlib.pyplot as plt

plt.plot(roc['FPR'], roc['TPR'], label='ROC Curve')
plt.xlabel('False Positive Rate (FPR)')
plt.ylabel('True Positive Rate (TPR)')
plt.title('ROC Curve')
plt.legend(loc='lower right')
print('Train AUC:', summary.areaUnderROC)
#plt.plot([0, 1], [0, 1], linestyle='--', color='gray') #randomclassifier
↳diagonal line
```

Train AUC: 0.8474526767588296



Do the same with pr from summary:

```
[175]: pr = summary.pr.toPandas()
pr
```

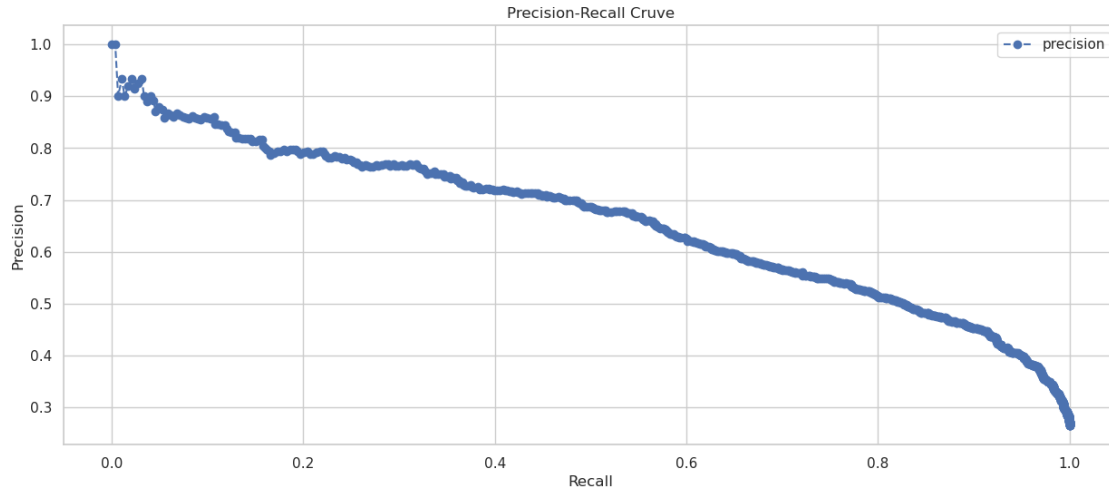
```
[175]:
```

	recall	precision
0	0.000000	1.000000
1	0.003754	1.000000
2	0.006757	0.900000
3	0.010511	0.933333
4	0.013514	0.900000
...
999	1.000000	0.265497
1000	1.000000	0.265233
1001	1.000000	0.264969
1002	1.000000	0.264706
1003	1.000000	0.264496

[1004 rows x 2 columns]

```
[176]: pr.plot(x='recall', y='precision', style='--o', legend=False)
plt.xlabel('Recall')
plt.title('Precision-Recall Cruve')
plt.ylabel('Precision')
plt.legend(loc='upper right')
```

```
[176]: <matplotlib.legend.Legend at 0x7f0485e25a20>
```



Our baseline model looks promising. Let's do some predictions on the `test` set.

Pass the `test` set through our trained model. Called the resulting DataFrame `fittedTest`:

```
[177]: fittedTest = lrModel.transform(test)
```

Print the first few rows of this DataFrame. Only show the following columns: "label", "prediction", "rawPrediction"

```
[178]: fittedTest.select('label', 'prediction', 'rawPrediction').show(6)
```

```
+-----+-----+-----+
|label|prediction|      rawPrediction|
+-----+-----+-----+
| 0.0|      1.0|[-0.3511581527595...|
| 1.0|      1.0|[-0.2767090139105...|
| 0.0|      0.0|[0.55682764262286...|
| 1.0|      1.0|[-0.2172075913226...|
| 0.0|      0.0|[0.51122949635904...|
| 1.0|      0.0|[0.04185624638468...|
+-----+-----+-----+
only showing top 6 rows
```

Make an evaluator from `BinaryClassificationEvaluator` function that calculates AUC. We will use this function to measure our model's performance on the `test` set. Call this evaluator `aucEvaluator`.

Note that this function can be found under the `pyspark.ml.evaluation` module.

```
[179]: from pyspark.ml.evaluation import BinaryClassificationEvaluator
from pyspark.mllib.evaluation import BinaryClassificationMetrics
aucEvaluator = BinaryClassificationEvaluator(metricName="areaUnderROC")
```

```
#fittedTest
```

Using our aucEvaluator find out the AUC on the test set:

```
[180]: auc_test = aucEvaluator.evaluate(fittedTest)
print(f"AUC on test set: {auc_test:.4f}")
```

AUC on test set: 0.8500

```
[181]: # fittedTestRDD = fittedTest.select("label", "prediction").rdd
# predictionAndLabels = fittedTestRDD.map(lambda row: (float(row.prediction),
# ↪row.label))
# metrics = BinaryClassificationMetrics(predictionAndLabels)
# auc = metrics.areaUnderROC
# print("Area under ROC = %s" % auc)
```

```
[182]: print('AUC on Train:', summary.areaUnderROC)
print(f"AUC on test set: {auc_test:.11f}")
```

AUC on Train: 0.8474526767588296

AUC on test set: 0.84998036459

Q. Are your test and train AUC's within the same range?

Ans. Yes both the AUC - Test & Train are within the same range. Overall, the results suggest that the logistic regression model is a decent classifier for this problem. However, there may still be room for improvement by experimenting with different models, feature engineering, or hyperparameter tuning.

1.4 2. Model Enhancement

- Explore different options to improve the logistic regression model using knowledge and examples from the previous three lectures.
- Implement at least two improvements to the model (e.g., feature selection, hyperparameter tuning, or additional feature engineering).

1.4.1 Feature Engineering

Creating number of opted services & MonthlyCharges/TotalCharges giving us the rate

```
[183]: services = ['PhoneService', 'MultipleLines', 'InternetService',
                  'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
                  'StreamingTV', 'StreamingMovies']
```

```
[184]: from pyspark.sql.functions import col, expr
num_of_services_expr = " + ".join([f"CAST(({c} = 'Yes') as INT)" for c in
# ↪services])
```

```
[185]: df = df.withColumn('numofservices', expr(num_of_services_expr))
```



```
[186]: #data = df.cache()
df = df.withColumn('charge_ratio', round(col('MonthlyCharges') /
↳ col('TotalCharges'), 2))
df.limit(2).toPandas()
```

```
[186]:      gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
0  Female                0      Yes          No         1           No
1   Male                0      No          No        34           Yes

      MultipleLines  InternetService  OnlineSecurity  OnlineBackup  \
0  NoPhoneService                DSL              No           Yes
1                No                DSL              Yes           No

      DeviceProtection  TechSupport  StreamingTV  StreamingMovies  Contract  \
0                No          No          No          No  Month_to_month
1                Yes          No          No          No    OneYear

      PaperlessBilling  PaymentMethod  MonthlyCharges  TotalCharges  Churn  \
0                Yes  ElectronicCheck         29.85         29.85    No
1                No    MailedCheck         56.95        1889.50    No

      numofservices  charge_ratio
0                1          1.00
1                3          0.03
```

```
[187]: [(c, df.where(col(c).isNull()).count()) for c in df.columns]
```

```
[187]: [('gender', 0),
      ('SeniorCitizen', 0),
      ('Partner', 0),
      ('Dependents', 0),
      ('tenure', 0),
      ('PhoneService', 0),
      ('MultipleLines', 0),
      ('InternetService', 0),
      ('OnlineSecurity', 0),
      ('OnlineBackup', 0),
      ('DeviceProtection', 0),
      ('TechSupport', 0),
      ('StreamingTV', 0),
      ('StreamingMovies', 0),
      ('Contract', 0),
      ('PaperlessBilling', 0),
      ('PaymentMethod', 0),
      ('MonthlyCharges', 0),
      ('TotalCharges', 0),
      ('Churn', 0),
```

```
('numofservices', 0),
('charge_ratio', 11)]
```

```
[188]: df = df.fillna(0)
```

```
[189]: rformula = RFormula(formula="Churn ~ . + numofservices + charge_ratio",
    featuresCol="features", labelCol="label")

fittedRF = rformula.fit(df)
preparedDF = fittedRF.transform(df)
train, test = preparedDF.randomSplit([0.7, 0.3], seed=42)
lrModel = lr.fit(train)

fittedTest = lrModel.transform(test)
testAUC = aucEvaluator.evaluate(fittedTest)

print(f"AUC on test set with new features: {testAUC:.4f}")
```

AUC on test set with new features: 0.8579

```
[190]: df.limit(10).toPandas()
```

```
[190]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	Female	0	Yes	No	1	No	
1	Male	0	No	No	34	Yes	
2	Male	0	No	No	2	Yes	
3	Male	0	No	No	45	No	
4	Female	0	No	No	2	Yes	
5	Female	0	No	No	8	Yes	
6	Male	0	No	Yes	22	Yes	
7	Female	0	No	No	10	No	
8	Female	0	Yes	No	28	Yes	
9	Male	0	No	Yes	62	Yes	

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	\
0	NoPhoneService	DSL	No	Yes	
1	No	DSL	Yes	No	
2	No	DSL	Yes	Yes	
3	NoPhoneService	DSL	Yes	No	
4	No	FiberOptic	No	No	
5	Yes	FiberOptic	No	No	
6	Yes	FiberOptic	No	Yes	
7	NoPhoneService	DSL	Yes	No	
8	Yes	FiberOptic	No	No	
9	No	DSL	Yes	Yes	

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
--	------------------	-------------	-------------	-----------------	----------	---

0	No	No	No	No	Month_to_month
1	Yes	No	No	No	OneYear
2	No	No	No	No	Month_to_month
3	Yes	Yes	No	No	OneYear
4	No	No	No	No	Month_to_month
5	Yes	No	Yes	Yes	Month_to_month
6	No	No	Yes	No	Month_to_month
7	No	No	No	No	Month_to_month
8	Yes	Yes	Yes	Yes	Month_to_month
9	No	No	No	No	OneYear

	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	Churn	\
0	Yes	ElectronicCheck	29.85	29.85	No	
1	No	MailedCheck	56.95	1889.50	No	
2	Yes	MailedCheck	53.85	108.15	Yes	
3	No	BankTransferAutomatic	42.30	1840.75	No	
4	Yes	ElectronicCheck	70.70	151.65	Yes	
5	Yes	ElectronicCheck	99.65	820.50	Yes	
6	Yes	CreditCardAutomatic	89.10	1949.40	No	
7	No	MailedCheck	29.75	301.90	No	
8	Yes	ElectronicCheck	104.80	3046.05	Yes	
9	No	BankTransferAutomatic	56.15	3487.95	No	

	numofservices	charge_ratio
0	1	1.00
1	3	0.03
2	3	0.50
3	3	0.02
4	1	0.47
5	5	0.12
6	4	0.05
7	1	0.10
8	6	0.03
9	3	0.02

1.4.2 Hyperparameter tuning using Grid Search and Cross-Validation

```
[191]: from pyspark.ml.tuning import CrossValidator, ParamGridBuilder
```

```
paramGrid = ParamGridBuilder() \
    .addGrid(lr.regParam, [0.0, 0.1, 0.01]) \
    .addGrid(lr.elasticNetParam, [0.0, 0.5, 1.0]) \
    .addGrid(lr.maxIter, [10, 50, 100]) \
    .build()

crossval = CrossValidator(estimator=lr,
                          estimatorParamMaps=paramGrid,
```

```

        evaluator=BinaryClassificationEvaluator(),
        numFolds=3)

cvModel = crossval.fit(train)

fittedTest = cvModel.transform(test)

auc_test = aucEvaluator.evaluate(fittedTest)

print(f"Test AUC: {auc_test:.4f}")

```

Test AUC: 0.8593

1.4.3 OneHotEncoding & RandomForestClassifier

```

[122]: from pyspark.ml.feature import StringIndexer, OneHotEncoder
       from pyspark.ml import Pipeline

dfv = df.cache()
categorical_columns = [
    "gender", "Partner", "Dependents", "PhoneService", "MultipleLines",
    ↪ "InternetService", "OnlineSecurity",
    "OnlineBackup", "DeviceProtection", "TechSupport", "StreamingTV",
    ↪ "StreamingMovies", "Contract",
    "PaperlessBilling", "PaymentMethod"
]

# Apply StringIndexer transformation
indexers = [StringIndexer(inputCol=column, outputCol=column+"_index").fit(df)
    ↪ for column in categorical_columns]
pipeline_indexers = Pipeline(stages=indexers)
df_indexed = pipeline_indexers.fit(dfv).transform(dfv)

churn_indexer = StringIndexer(inputCol="Churn", outputCol="Churn_index")
df_transformed = churn_indexer.fit(df_transformed).transform(df_transformed)

# Apply OneHotEncoder transformation
encoders = [OneHotEncoder(inputCol=column+"_index", outputCol=column+"_Vec")
    ↪ for column in categorical_columns]
pipeline_encoders = Pipeline(stages=encoders)
df_transformed = pipeline_encoders.fit(df_indexed).transform(df_indexed)

```

23/04/26 14:14:43 WARN CacheManager: Asked to cache already cached data.

```

[124]: churn_indexer = StringIndexer(inputCol="Churn", outputCol="Churn_index")
       df_transformed = churn_indexer.fit(df_transformed).transform(df_transformed)

```

```
[125]: df_transformed.limit(9).toPandas()
```

```
[125]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	Female	0	Yes	No	1	No	
1	Male	0	No	No	34	Yes	
2	Male	0	No	No	2	Yes	
3	Male	0	No	No	45	No	
4	Female	0	No	No	2	Yes	
5	Female	0	No	No	8	Yes	
6	Male	0	No	Yes	22	Yes	
7	Female	0	No	No	10	No	
8	Female	0	Yes	No	28	Yes	

	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	\
0	NoPhoneService	DSL	No	Yes	
1	No	DSL	Yes	No	
2	No	DSL	Yes	Yes	
3	NoPhoneService	DSL	Yes	No	
4	No	FiberOptic	No	No	
5	Yes	FiberOptic	No	No	
6	Yes	FiberOptic	No	Yes	
7	NoPhoneService	DSL	Yes	No	
8	Yes	FiberOptic	No	No	

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
0	No	No	No	No	Month_to_month	
1	Yes	No	No	No	OneYear	
2	No	No	No	No	Month_to_month	
3	Yes	Yes	No	No	OneYear	
4	No	No	No	No	Month_to_month	
5	Yes	No	Yes	Yes	Month_to_month	
6	No	No	Yes	No	Month_to_month	
7	No	No	No	No	Month_to_month	
8	Yes	Yes	Yes	Yes	Month_to_month	

	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	Churn	\
0	Yes	ElectronicCheck	29.85	29.85	No	
1	No	MailedCheck	56.95	1889.50	No	
2	Yes	MailedCheck	53.85	108.15	Yes	
3	No	BankTransferAutomatic	42.30	1840.75	No	
4	Yes	ElectronicCheck	70.70	151.65	Yes	
5	Yes	ElectronicCheck	99.65	820.50	Yes	
6	Yes	CreditCardAutomatic	89.10	1949.40	No	
7	No	MailedCheck	29.75	301.90	No	
8	Yes	ElectronicCheck	104.80	3046.05	Yes	

	numofservices	charge_ratio	gender_index	Partner_index	Dependents_index	\
--	---------------	--------------	--------------	---------------	------------------	---

0	1	1.00	1.0	1.0	0.0
1	3	0.03	0.0	0.0	0.0
2	3	0.50	0.0	0.0	0.0
3	3	0.02	0.0	0.0	0.0
4	1	0.47	1.0	0.0	0.0
5	5	0.12	1.0	0.0	0.0
6	4	0.05	0.0	0.0	1.0
7	1	0.10	1.0	0.0	0.0
8	6	0.03	1.0	1.0	0.0

	PhoneService_index	MultipleLines_index	InternetService_index	\
0	1.0	2.0	1.0	
1	0.0	0.0	1.0	
2	0.0	0.0	1.0	
3	1.0	2.0	1.0	
4	0.0	0.0	0.0	
5	0.0	1.0	0.0	
6	0.0	1.0	0.0	
7	1.0	2.0	1.0	
8	0.0	1.0	0.0	

	OnlineSecurity_index	OnlineBackup_index	DeviceProtection_index	\
0	0.0	1.0	0.0	
1	1.0	0.0	1.0	
2	1.0	1.0	0.0	
3	1.0	0.0	1.0	
4	0.0	0.0	0.0	
5	0.0	0.0	1.0	
6	0.0	1.0	0.0	
7	1.0	0.0	0.0	
8	0.0	0.0	1.0	

	TechSupport_index	StreamingTV_index	StreamingMovies_index	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	1.0	0.0	0.0	
4	0.0	0.0	0.0	
5	0.0	1.0	1.0	
6	0.0	1.0	0.0	
7	0.0	0.0	0.0	
8	1.0	1.0	1.0	

	Contract_index	PaperlessBilling_index	PaymentMethod_index	gender_Vec	\
0	0.0	0.0	0.0	(0.0)	
1	2.0	1.0	1.0	(1.0)	
2	0.0	0.0	1.0	(1.0)	

3	2.0	1.0	2.0	(1.0)
4	0.0	0.0	0.0	(0.0)
5	0.0	0.0	0.0	(0.0)
6	0.0	0.0	3.0	(1.0)
7	0.0	1.0	1.0	(0.0)
8	0.0	0.0	0.0	(0.0)

	Partner_Vec	Dependents_Vec	PhoneService_Vec	MultipleLines_Vec	\
0	(0.0)	(1.0)	(0.0)	(0.0, 0.0)	
1	(1.0)	(1.0)	(1.0)	(1.0, 0.0)	
2	(1.0)	(1.0)	(1.0)	(1.0, 0.0)	
3	(1.0)	(1.0)	(0.0)	(0.0, 0.0)	
4	(1.0)	(1.0)	(1.0)	(1.0, 0.0)	
5	(1.0)	(1.0)	(1.0)	(0.0, 1.0)	
6	(1.0)	(0.0)	(1.0)	(0.0, 1.0)	
7	(1.0)	(1.0)	(0.0)	(0.0, 0.0)	
8	(0.0)	(1.0)	(1.0)	(0.0, 1.0)	

	InternetService_Vec	OnlineSecurity_Vec	OnlineBackup_Vec	\
0	(0.0, 1.0)	(1.0, 0.0)	(0.0, 1.0)	
1	(0.0, 1.0)	(0.0, 1.0)	(1.0, 0.0)	
2	(0.0, 1.0)	(0.0, 1.0)	(0.0, 1.0)	
3	(0.0, 1.0)	(0.0, 1.0)	(1.0, 0.0)	
4	(1.0, 0.0)	(1.0, 0.0)	(1.0, 0.0)	
5	(1.0, 0.0)	(1.0, 0.0)	(1.0, 0.0)	
6	(1.0, 0.0)	(1.0, 0.0)	(0.0, 1.0)	
7	(0.0, 1.0)	(0.0, 1.0)	(1.0, 0.0)	
8	(1.0, 0.0)	(1.0, 0.0)	(1.0, 0.0)	

	DeviceProtection_Vec	TechSupport_Vec	StreamingTV_Vec	StreamingMovies_Vec	\
0	(1.0, 0.0)	(1.0, 0.0)	(1.0, 0.0)	(1.0, 0.0)	
1	(0.0, 1.0)	(1.0, 0.0)	(1.0, 0.0)	(1.0, 0.0)	
2	(1.0, 0.0)	(1.0, 0.0)	(1.0, 0.0)	(1.0, 0.0)	
3	(0.0, 1.0)	(0.0, 1.0)	(1.0, 0.0)	(1.0, 0.0)	
4	(1.0, 0.0)	(1.0, 0.0)	(1.0, 0.0)	(1.0, 0.0)	
5	(0.0, 1.0)	(1.0, 0.0)	(0.0, 1.0)	(0.0, 1.0)	
6	(1.0, 0.0)	(1.0, 0.0)	(0.0, 1.0)	(1.0, 0.0)	
7	(1.0, 0.0)	(1.0, 0.0)	(1.0, 0.0)	(1.0, 0.0)	
8	(0.0, 1.0)	(0.0, 1.0)	(0.0, 1.0)	(0.0, 1.0)	

	Contract_Vec	PaperlessBilling_Vec	PaymentMethod_Vec	Churn_index
0	(1.0, 0.0)	(1.0)	(1.0, 0.0, 0.0)	0.0
1	(0.0, 0.0)	(0.0)	(0.0, 1.0, 0.0)	0.0
2	(1.0, 0.0)	(1.0)	(0.0, 1.0, 0.0)	1.0
3	(0.0, 0.0)	(0.0)	(0.0, 0.0, 1.0)	0.0
4	(1.0, 0.0)	(1.0)	(1.0, 0.0, 0.0)	1.0
5	(1.0, 0.0)	(1.0)	(1.0, 0.0, 0.0)	1.0

6	(1.0, 0.0)	(1.0)	(0.0, 0.0, 0.0)	0.0
7	(1.0, 0.0)	(0.0)	(0.0, 1.0, 0.0)	0.0
8	(1.0, 0.0)	(1.0)	(1.0, 0.0, 0.0)	1.0

```
[127]: from pyspark.ml.feature import VectorAssembler
from pyspark.ml.classification import RandomForestClassifier
from pyspark.ml.evaluation import BinaryClassificationEvaluator

assembler = VectorAssembler(
    inputCols=[column+"_Vec" for column in categorical_columns] + ["tenure",
↳"MonthlyCharges", "TotalCharges", "numofservices", "charge_ratio"],
    outputCol="features"
)
df_assembled = assembler.transform(df_transformed)

train, test = df_assembled.randomSplit([0.8, 0.2], seed=42)
rf = RandomForestClassifier(labelCol="Churn_index", featuresCol="features",
↳numTrees=100)
model = rf.fit(train)

predictions = model.transform(test)
evaluator = BinaryClassificationEvaluator(labelCol="Churn_index",
↳rawPredictionCol="rawPrediction", metricName="areaUnderROC")
auc = evaluator.evaluate(predictions)
print("Area Under ROC Curve: ", auc)
```

Area Under ROC Curve: 0.8440139671275291

1.5 3. Next Steps

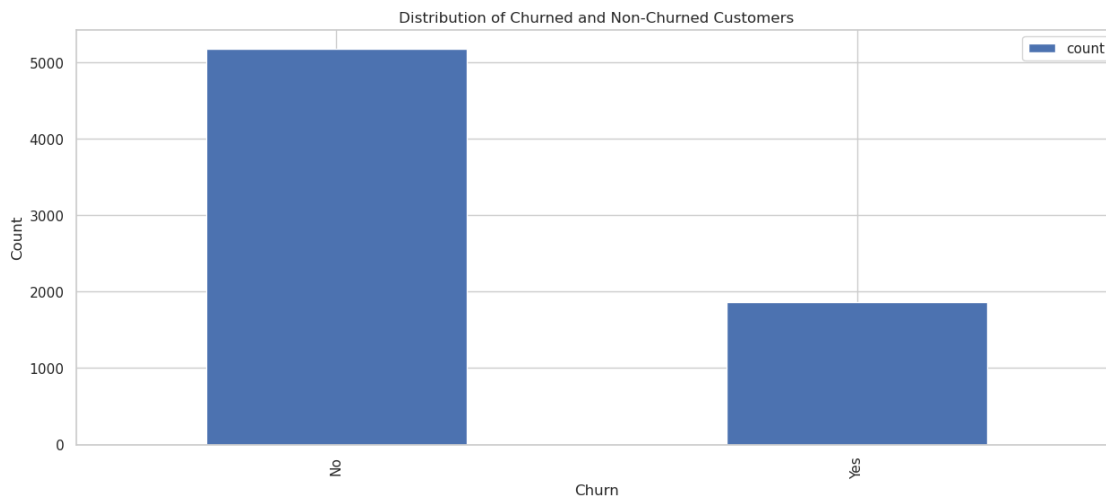
Provide a list of additional suggestions to further enhance the model or feature engineering techniques to try.

Discuss the potential impact of these suggestions on the model's performance.

1. Feature scaling: We can perform scaling on the numeric columns using StandardScaler on ['tenure', 'MonthlyCharges', 'TotalCharges'], resulting in similar range of values and improved performance of the models.
2. Feature Engineering: We can create a new column for overall revenue. This can be done by multiplying tenure with monthlycharges. 2.1 As a feature engineering technique, we can create a new column to calculate the overall revenue by multiplying tenure with monthlycharges. Additionally, we can perform an in-depth analysis of the relationship between tenure and contract to gain insights into customer churn. For example, we can explore the correlation between 'year to year contract' and 'tenure < 10 months' to identify customers who are less likely to churn. Conversely, we can focus on customers who have been on a 'year to year' contract & tenure > 10 months or more as they may be more likely to churn.

3. Feature selection: We can also use techniques such as correlation analysis, forward/backward selection, and LASSO regularization to identify and remove redundant or irrelevant features that may negatively impact the model's performance.
4. Resampling Techniques: From the below graph we can see that the dataset is imbalanced, having majority as not churning (NO > 5000). Therefore, oversampling the minority and undersampling the majority can help improve the model's performance. (Note: Oversampling -> can also lead to -> Overfitting & Undersampling -> to -> Underfitting)

```
[142]: # 4. Point
churn_count = df.groupby('Churn').count().toPandas()
churn_count.plot(kind='bar', x='Churn', y='count')
plt.title('Distribution of Churned and Non-Churned Customers')
plt.xlabel('Churn')
plt.ylabel('Count')
plt.show()
```



1.6 4. Final Model Evaluation

Report your best model with its AUC on the test and train sets.

```
[192]: # Best Model with AUC on train and test sets.
train_auc = aucEvaluator.evaluate(cvModel.transform(train)) #train
print(f"Train AUC: {train_auc:.4f}")

# AUC on test set
test_auc = aucEvaluator.evaluate(cvModel.transform(test)) #test
print(f"Test AUC: {test_auc:.4f}")
```

Train AUC: 0.8498

Test AUC: 0.8593

From the above the best model achieved so far is from hyper-parameter tuning having AUC (train): 0.8498 and AUC (test): 0.8593 - explaining the model's ability to interpret positive and negative instances.

1.7 5. Notebook Organization

References: 1. <https://github.com/soltaniehha/Big-Data-Analytics-for-Business> : for lecture notes & .ipynb notebook and previous lecture concept implementation 2. https://www.w3schools.com/python/matplotlib_pyplot.asp : Used for better graph learning opportunity and practice. 3. notebook 11-Classification-Regression/02-Logistic-Regression-Example.ipynb

```
[204]: # bm = cvModel.bestModel.summary
# print(bm)
# bestModelroc = bm.roc.toPandas()
# bestModelroc
# %matplotlib inline
# import matplotlib.pyplot as plt

# plt.plot(bestModelroc['FPR'], bestModelroc['TPR'], label='ROC Curve')
# plt.xlabel('False Positive Rate (FPR)')
# plt.ylabel('True Positive Rate (TPR)')
# plt.title('ROC Curve')
# plt.legend(loc='lower right')
# #print('Train AUC:', summary.areaUnderROC)
# bestModelpr = bm.pr.toPandas()
# bestModelpr
# bestModelpr.plot(x='recall', y='precision', style='--o', legend=False)
# plt.xlabel('Recall')
# plt.title('Precision-Recall Cruve')
# plt.ylabel('Precision')
# plt.legend(loc='upper right')
```